

# MineSweeper Solver

March 15, 2019

## 1 Training a Convolutional Neural Network to solve minesweeper

### 1.1 Description

I wanted to learn a bit of tensorflow, and also to solve the minesweeper problem. At first, (a year ago), I thought to use basic machine learning, but didn't succeed because I barely knew how to do anything. The idea was to train a network, that gets a partially open board and tries to predict where the mines are on the board. We then use this network to look at the game-board and find the square with the least likelihood of being a mine. We then open it and predict once more on the updated game-board, and so on. `## Code`

```
In [2]: ### Importing libraries
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import os

In [3]: ### Create board
def createBoard(boardSize, mineNum):
    """
        Creates a board given the size of the board (a square), and the number of
        mines. It places the mines randomly, then runs over all squares and
        whenever there is a mine, it adds 1 to all 8 neighbours.
    """
    # put the mines in random places
    A = np.zeros([boardSize, boardSize])
    for ii in range(mineNum): # randomly allocate mine
        randX = np.random.randint(0, boardSize - 1)
        randY = np.random.randint(0, boardSize - 1)
        while A[randX, randY] == 9: # while the place already has a mine
            randX = np.random.randint(0, boardSize - 1)
            randY = np.random.randint(0, boardSize - 1)
        A[randX, randY] = 9
    B = np.zeros([boardSize, boardSize])
    midY = boardSize / 2 #-(1-size)/2
    midX = midY
```

```

# the following creates the numbers around the board
for y1 in range(boardSize):
    for x1 in range(boardSize):
        if A[y1, x1]==9:
            for e in [-1, 0, 1]:
                for f in [-1, 0, 1]:
                    if e*e + f*f != 0 and abs(y1 + e - midY) <= midY and \
                        abs(x1 + f - midX) <= midX:
                        if A[y1 + e, x1 + f] != 9:
                            B[y1 + e, x1 + f] = B[y1 + e, x1 + f] + 1
board = A + B
return board

```

In [4]: *%% Draw board*

```

def drawBoard(boardMat, imgPath):
    """
        Receives the matrix description of the board, and the path to the images
        of the numbers and mine and flag, and creates the board.
    """
    # load images
    currImg = mpimg.imread(imgPath + '\\\' + str(0) + '.png')
    imshp = currImg.shape[:2] # relies on the fact that all images are the same size
    boardImg = np.zeros([boardMat.shape[0] * imshp[0], boardMat.shape[1] * imshp[1], 3])
    for ii in range(boardMat.shape[0]):
        for jj in range(boardMat.shape[1]):
            boardImg[ii*imshp[0]:(ii+1)*imshp[0], jj*imshp[1]:(jj+1)*imshp[1], :] = \
                mpimg.imread(imgPath + '\\\' + str(int(boardMat[ii, jj])) + '.png')
    plt.figure(figsize=(8, 8))
    plt.imshow(boardImg)
    return boardImg

```

In [5]: *%% Create partial board for training, and testing*

```

def createPartialBoards(N, boardSize, mineNum): # N is the number of boards
    """ We'll get a random number of open squares for each board, create N boards,
        and then get a partial solution and the binary solution"""
    boardGT = [] # list of boards
    openRand = np.random.randint(boardSize*boardSize, size = N) # random number of open
    openRand = openRand.tolist()
    trainBrd = [] # list of the partial boards for training
    boardSol = [] # list of the solutions for the training boards
    for ii in range(N):
        # create board
        currBoard = createBoard(boardSize, mineNum)
        boardGT.append(currBoard)
        # create partial board
        currRand = np.random.randint(boardSize, size = [openRand[ii], 2]) # the location of
        currRand = currRand.tolist()
        currPrtl = 10 * np.ones([boardSize, boardSize])

```

```

    for jj in range(openRand[ii]):
        currPrt1[currRand[jj][0], currRand[jj][1]] = currBoard[currRand[jj][0], currRand[jj][1]]
    trainBrd.append(currPrt1)
    # create solution
    currSol = (currBoard == 9)
    boardSol.append(currSol)
boardGT = np.array(boardGT)
boardGT = boardGT[:,:,:,:np.newaxis]
trainBrd = np.array(trainBrd)
trainBrd = trainBrd[:,:,:,:np.newaxis]
boardSol = np.array(boardSol)
boardSol = boardSol[:,:,:,:np.newaxis]
return boardGT, trainBrd, boardSol

```

In [6]: *%% Create network*

```

def createCNN(boardSize):
    CNet = tf.keras.Sequential()
    CNet.add(tf.keras.layers.Conv2D(32, 3, activation = 'relu',
        input_shape = (boardSize, boardSize, 1), name = 'conv1'))
    CNet.add(tf.keras.layers.Conv2D(64, 3, activation = 'relu', name = 'conv2')) # use_b
    CNet.add(tf.keras.layers.Conv2D(128, 3, activation = 'relu', name = 'conv3'))
    CNet.add(tf.keras.layers.Conv2DTranspose(128, 3, activation = 'relu', name = 'deconv1'))
    CNet.add(tf.keras.layers.Conv2DTranspose(128, 3, activation = 'relu', name = 'deconv2'))
    CNet.add(tf.keras.layers.Conv2DTranspose(64, 3, activation = 'relu', name = 'deconv3'))
    #CNet.add(tf.keras.layers.Dense(32, name = 'dense1'))
    CNet.add(tf.keras.layers.Dense(1, name = 'dense1'))
    # compile
    CNet.compile(optimizer = tf.train.AdamOptimizer(0.001), loss = 'mse', metrics = ['acc'])
    print('CNet.summary():')
    CNet.summary()
    return CNet

```

In [7]: *%% Show prediction output*

```

def showPrediction(currGT, currBrd, currSol, currOut, thresh, imgPath):
    GTImg = drawBoard(currGT, imgPath)
    BrdImg = drawBoard(currBrd, imgPath)
    SolImg = drawBoard(9*currSol, imgPath)
    # Out is between 0 and 1, so we'll use a threshold, and then multiply by 9:
    outMat = 9*(currOut > thresh)
    OutImg = drawBoard(outMat, imgPath)
    plt.figure(figsize=(15,15))
    plt.subplot(221)
    plt.imshow(GTImg)
    plt.title('Ground Truth')
    plt.subplot(222)
    plt.imshow(BrdImg)
    plt.title('Training board')
    plt.subplot(223)
    plt.imshow(OutImg)
    plt.title('Prediction')

```

```

plt.imshow(SolImg)
plt.title('Real solution')
plt.subplot(224)
plt.imshow(OutImg)
plt.title('Auto solution')

```

In [8]: *%% open zeros neighbors opens all the neighbors of some index, given that its # value is 0.*

```

def openZerosNeighbors(gameBoard, GTBoard, ind, cntr):
    midY = gameBoard.shape[1] / 2 #-(1-size)/2
    midX = midY
    Y = gameBoard.shape[1]
    X = gameBoard.shape[1]
    for e in [-1, 0, 1]:
        for f in [-1, 0, 1]:
            newInd = [0, ind[1] + e, ind[2] + f, 0]
            # print(newInd)
            if e*e + f*f > 0 and newInd[1] >= 0 and newInd[1] < Y and newInd[2] >= 0 and newInd[2] < X:
                #if e*e + f*f > 0 and abs(newInd[1] - midY) <= midY and abs(newInd[2] - midX) <= midX:
                if gameBoard[newInd[0], newInd[1], newInd[2], newInd[3]] == 10:
                    gameBoard[newInd[0], newInd[1], newInd[2], newInd[3]] = GTBoard[newInd[1], newInd[2]]
                    cntr += 1
                if gameBoard[newInd[0], newInd[1], newInd[2], newInd[3]] == 0: # we landed on a mine
                    gameBoard, cntr = openZerosNeighbors(gameBoard, GTBoard, newInd, cntr)
    return gameBoard, cntr

```

In [9]: *%% Play games and check if you win or lose, and return the number of steps until you win or lose.*

```

def playGame(boardSize, mineNum, CNet, imgPath, showOutcome, open0, expFlag):
    # showOutcome is a boolean for whether to print and show or not
    GTBoard = createBoard(boardSize, mineNum)
    # We also need the changing board:
    gameBoard = 10 * np.ones([1, boardSize, boardSize, 1])
    # And now we need a while loop that predicts the board and opens the least probable square (closest to nil):
    done = 0
    cntr = 0 # counter of the number of steps
    while done == 0:
        cntr = cntr + 1
        if expFlag:
            predBoard = CNet.predict(np.log10(gameBoard + 0.1))
        else:
            predBoard = CNet.predict(gameBoard)
        # We need to force the machine to pick squares he hasn't picked yet, so we subtract the predicted board from the game board:
        ind = np.unravel_index(np.argmin(predBoard - gameBoard, axis=None), predBoard.shape)
        gameBoard[ind] = GTBoard[ind[1:3]]
        if GTBoard[ind[1:3]] == 9: # landed on a mine
            done = 1
        if showOutcome:

```

```

        print('You LOST :(')
    if cntr == boardSize**2 - mineNum:
        done = 1
        if showOutcome:
            print('You WON :)')
    if open0 and gameBoard[ind] == 0: # open0 - open neighbors of 0, and we landed on
        gameBoard, cntr = openZerosNeighbors(gameBoard, GTBoard, ind, cntr)
# Show the outcome
if showOutcome:
    GTImg = drawBoard(GTBoard, imgPath)
    gameImg = drawBoard(np.squeeze(gameBoard), imgPath)
    plt.figure(figsize=(16, 16))
    plt.subplot(211)
    plt.imshow(GTImg)
    plt.title('Ground Truth')
    plt.subplot(212)
    plt.imshow(gameImg)
    plt.title('Game Board')
    print(cntr)
return cntr

```

In [10]: *%% test the program:*

```

def testCNet(testN, showOutcome, open0, boardSize, mineNum, imgPath, expFlag):
    maxCnt = boardSize**2 - mineNum # the number of steps needed to win
    stepCnt = np.zeros(testN)
    for ii in range(testN):
        stepCnt[ii] = playGame(boardSize, mineNum, CNet, imgPath, showOutcome, open0, expFlag)

    # show the histogram:
    plt.figure(figsize=(10,10))
    gamehist = plt.hist(stepCnt, maxCnt)
    winRate = np.sum(stepCnt == maxCnt) / (testN - np.sum(stepCnt == 1)) # success rate
    print(winRate)
    return 0


```

In [12]: *%% Main:*

```

# create a board as an example, and then show it:
# We'll try intermediate, which is 16x16 with 40 mines
boardSize = 16
mineNum = 40
board = createBoard(boardSize, mineNum)
imgPath = './\MineSweeper' # the path where the images of the
boardImg = drawBoard(board, imgPath) # draw the board in an understandable manner

```



```
In [13]: ### create data:
# We want partial boards in all the different stages, and their solutions - a
# binary board with 1's for mines and 0's for all else.
# Lets assume we have N boards:
N = 400000
expFlag = 1
# Now we want to create the boards:
boardDat, trainBrd, boardSol = createPartialBoards(N, boardSize, mineNum)
if expFlag:
    boardDat = np.log10(boardDat + 0.1)
    trainBrd = np.log10(trainBrd + 0.1)
# split into training and validation
splitInt = int(0.9 * N)
validDat = trainBrd[splitInt:]
```

```

validSol = boardSol[splitInt:]
trainDat = trainBrd[:splitInt]
trainSol = boardSol[:splitInt]

```

In [14]: *### Create Network*

*# Saving network*

```
checkpoint_path = ".\\checkpoint9\\cp.ckpt"
```

```
checkpoint_dir = os.path.dirname(checkpoint_path)
```

```
cp_callback = tf.keras.callbacks.ModelCheckpoint(checkpoint_path,
                                                  save_weights_only=True,
                                                  verbose=1)
```

*# create the network*

```
CNet = createCNN(boardSize)
```

*# Now we want to train the network:*

```
CNet.fit(trainDat, trainSol, epochs = 10, batch_size = 128,
        validation_data = (validDat, validSol), callbacks = [cp_callback])
```

WARNING:tensorflow:From C:\Anaconda\envs\tf-cpu\lib\site-packages\tensorflow\python\ops\resource\_ops.py:111: tf.nn.conv2d is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

WARNING:tensorflow:From C:\Anaconda\envs\tf-cpu\lib\site-packages\tensorflow\python\keras\utils\tf\_utils.py:111: tf.nn.conv2d is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.nn.conv2d instead.

CNet.summary():

Layer (type)	Output Shape	Param #
conv1 (Conv2D)	(None, 14, 14, 32)	320
conv2 (Conv2D)	(None, 12, 12, 64)	18496
conv3 (Conv2D)	(None, 10, 10, 128)	73856
deconv1 (Conv2DTranspose)	(None, 12, 12, 128)	147584
deconv2 (Conv2DTranspose)	(None, 14, 14, 128)	147584
deconv3 (Conv2DTranspose)	(None, 16, 16, 64)	73792
dense1 (Dense)	(None, 16, 16, 1)	65

Total params: 461,697

Trainable params: 461,697

Non-trainable params: 0

Train on 360000 samples, validate on 40000 samples

WARNING:tensorflow:From C:\Anaconda\envs\tf-cpu\lib\site-packages\tensorflow\python\ops\math\_ops.py:111: tf.nn.conv2d is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Epoch 1/10

359936/360000 [=====>.] - ETA: 0s - loss: 0.0630 - acc: 0.9140

Epoch 00001: saving model to .\checkpoint9\cp.ckpt

WARNING:tensorflow:From C:\Anaconda\envs\tf-cpu\lib\site-packages\tensorflow\python\keras\engi

Instructions for updating:

Use tf.train.CheckpointManager to manage checkpoints rather than manually editing the Checkpoi

360000/360000 [=====] - 1740s 5ms/sample - loss: 0.0630 - acc: 0.9140

Epoch 2/10

359936/360000 [=====>.] - ETA: 0s - loss: 0.0490 - acc: 0.9339

Epoch 00002: saving model to .\checkpoint9\cp.ckpt

360000/360000 [=====] - 1637s 5ms/sample - loss: 0.0490 - acc: 0.9339

Epoch 3/10

359936/360000 [=====>.] - ETA: 0s - loss: 0.0459 - acc: 0.9384

Epoch 00003: saving model to .\checkpoint9\cp.ckpt

360000/360000 [=====] - 1639s 5ms/sample - loss: 0.0459 - acc: 0.9384

Epoch 4/10

359936/360000 [=====>.] - ETA: 0s - loss: 0.0445 - acc: 0.9402

Epoch 00004: saving model to .\checkpoint9\cp.ckpt

360000/360000 [=====] - 1642s 5ms/sample - loss: 0.0445 - acc: 0.9402

Epoch 5/10

359936/360000 [=====>.] - ETA: 0s - loss: 0.0437 - acc: 0.9413

Epoch 00005: saving model to .\checkpoint9\cp.ckpt

360000/360000 [=====] - 1638s 5ms/sample - loss: 0.0437 - acc: 0.9413

Epoch 6/10

359936/360000 [=====>.] - ETA: 0s - loss: 0.0431 - acc: 0.9420

Epoch 00006: saving model to .\checkpoint9\cp.ckpt

360000/360000 [=====] - 1637s 5ms/sample - loss: 0.0431 - acc: 0.9420

Epoch 7/10

359936/360000 [=====>.] - ETA: 0s - loss: 0.0427 - acc: 0.9425

Epoch 00007: saving model to .\checkpoint9\cp.ckpt

360000/360000 [=====] - 1636s 5ms/sample - loss: 0.0427 - acc: 0.9425

Epoch 8/10

359936/360000 [=====>.] - ETA: 0s - loss: 0.0424 - acc: 0.9429

Epoch 00008: saving model to .\checkpoint9\cp.ckpt

360000/360000 [=====] - 1639s 5ms/sample - loss: 0.0424 - acc: 0.9429

Epoch 9/10

359936/360000 [=====>.] - ETA: 0s - loss: 0.0422 - acc: 0.9432

Epoch 00009: saving model to .\checkpoint9\cp.ckpt

360000/360000 [=====] - 1637s 5ms/sample - loss: 0.0422 - acc: 0.9432

Epoch 10/10

359936/360000 [=====>.] - ETA: 0s - loss: 0.0420 - acc: 0.9434

Epoch 00010: saving model to .\checkpoint9\cp.ckpt

360000/360000 [=====] - 1675s 5ms/sample - loss: 0.0420 - acc: 0.9434

Out[14]: <tensorflow.python.keras.callbacks.History at 0x119cda694a8>


In [15]: #%% Test:




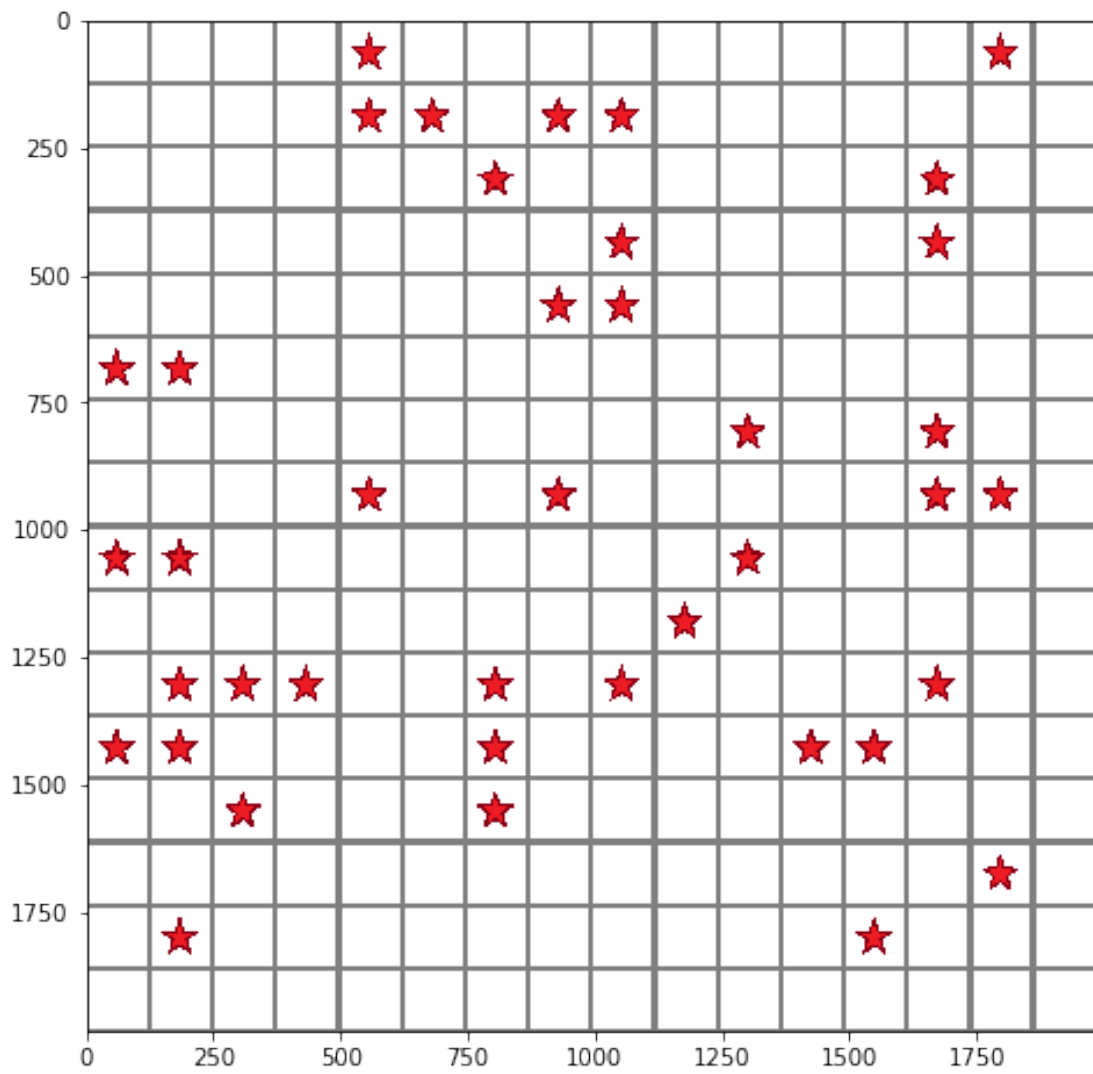
```

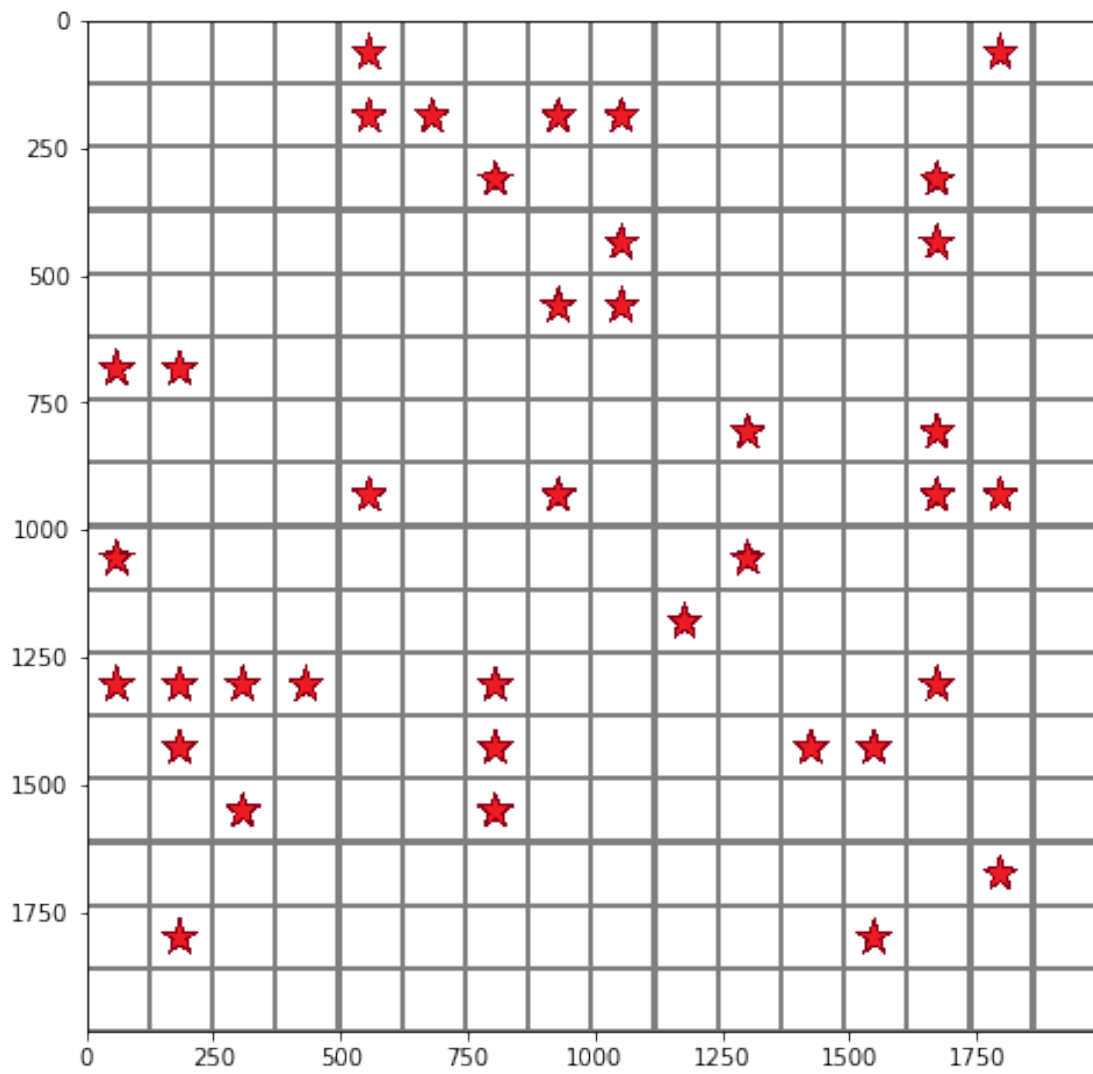
# create another board, and look at the solution
threshold = 0.5 # threshold above which we think we have a mine
#create a board:
testGT, testBrd, testSol = createPartialBoards(1, boardSize, mineNum)
testOut = CNet.predict(np.log10(testBrd+0.1))
ii = 0;
currGT = np.squeeze(testGT [ii,:,:])
currBrd = np.squeeze(testBrd[ii,:,:])
currSol = np.squeeze(testSol[ii,:,:])
currOut = np.squeeze(testOut[ii,:,:])
showPrediction(currGT, currBrd, currSol, currOut, threshold, imgPath)





```












```
In [16]: ### Now we want a system that gets a board, and plays it.
# We'll make a game board, and then use it to
testN = 1000 # number of games to test
showOutcome = 0
open0 = 1
test = testCNet(testN, showOutcome, open0, boardSize, mineNum, imgPath, expFlag)
```

0.662



## 1.2 I tried various networks: